

DFT RESEARCH BRIEF — Elite Squads & Sucky Squads

Tier definitions, the GS case, and the killer-leader inversion · Jul 16, 2026 · data: 2024–2026 WNBA, 1,220 team-games, 5,260 checkpoints · commits e229c8b + 2f2d959

Executive summary

1. The production elite cut (current record $>.600$) is a fine *gate* but a bad *description*: it uses future information and rewrites history (104 retroactive label flips in 2026 alone). Manny's definition — **ever reached .600 with 15+ games played** — is adopted as the descriptive lens: earned membership, one-way updates, best empirical ordering. Under it, GS transforms from 1-5 vs top into **8-6, the league's 2nd-best elite record**, and grades as its only DUAL_EDGE team.
2. No definition's vs-top split reliably *predicts* elite-vs-elite winners (LOW power, one season reverses). Raw record stays the only gap input. Production gates untouched.
3. For sucky squads, the production $<.400$ fade cut holds. But the proposed "exclude elite-killers" refinement **inverts**: bad teams with elite scalps are the *best* leads to fade — **64.7% trailer conversion inside the system band**, the strongest cell in the study — because a bad team that occasionally beats elites is a variance team, and variance teams don't hold leads against structure.
4. The killer flag is fully systematic — computable as-of at push time from data the nightly already stores — and composes with the existing quality-gap machinery. Proposed: **KILLER_LEADER ledger stream** (log-only, predicted ~65%), graduating exactly like GAP_BASE did.

PART I — ELITE SQUADS: what is a top team?

The problem, in one game

GS beat IND on Jul 15. IND fell to .583 — and under the production cut, GS's two wins over IND were retroactively relabeled "vs rest." GS beat a .600+ team both times *at the time of play*; the ledger now denies it. Same mechanism poisoned the NY (entered 12-7) and ATL (12-8) reads. A definition that rewrites the past is measuring the boundary, not the team.




Fig 1 — The league sorted honestly. Elite tier (cyan) = ever .600 @ 15+ GP. Distance above the amber line = record inflation. MIN and NY sit nearest the line (records earned); DAL/TOR/GS carry the biggest weak-slate cushions. Manny's eye-test distrust of ATL/IND/NY vs the true four shows up as the vs-rest gap: **MIN/GS/LV/DAL are a combined 44-7 (.863) against the rest of the league.**

Same games, three definitions — the GS "paper tiger" was an artifact




Fig 2 — The motivating case. GS under the production cut looked like the league's worst elite (1-5, -6.8 eFG). Under earned membership they are 8-6 at dead-even shooting, and 8-1 (+4.6 eFG) against the .575 band. The honest full sentence: **GS is an elite team with a super-elite problem** — 0-5 vs the MIN/LV core at -8.5 eFG with opponents shooting 44% from three. That part survives every definition.




Fig 3 — The predictive test (walk-forward, both teams $\geq .550$ as-of). EVER600 leads pooled and in 2 of 3 seasons, but 2025 reverses everything and per-season n is 23–54. Verdict: **ordering only — no vs-top split earns a production role**. This matches the honest-gap rejection: raw record remains the sole gap input.

Insights — elite

- I-1. "Elite" should be earned and sticky: 15 games of proof to enter, no exit on a slump, no entry on a hot start. Adopted as the descriptive lens for team ctx, briefs, and reads like GS.
- I-2. The four true elites going 44-7 vs rest is the lane logic's foundation — but it is a retrospective, heavily market-priced number. Its betting expression is the OVERRIDE LANE (negative-inflation top teams at soft live prices), not blind favorite backing.
- I-3. Redefinition is not free: under EVER600 every team's inflation compresses and MIN's flips positive — the override lane would empty. Lane and registration stay on def A numbers.

PART II — SUCKY SQUADS: what is a beatable leader?

Frame: the locked-signal substrate — 5,260 historical checkpoints → Q2/Q3 spots, trailer down 1–9, first checkpoint per game-quarter, both teams ≥ 4 GP, strictly as-of records. **n=744** (2024: 297, 2025: 447), baseline trailer win 34.1%. Seven definitions pre-registered before running.




Fig 4 — Every "bad leader" definition roughly doubles trailer conversion vs non-qualifying spots (~29%). Margin/game ≤ -4 (green) is the best and most season-stable standalone (49/54) — badness is better measured continuously than by record membership. The production $<.400$ cut is statistically indistinguishable from it: **no gate change warranted**. The red bar is the story →




Fig 5 — The inversion. Excluding elite-killers (the hypothesis) produced the WORST definition tested; isolating them produced the best cell in the study: 58.2% standalone, **64.7% inside the band (n=51, 60%/69% by season)** — nine points over the band base. Why it coheres with the locked thesis: a sub-.450 team with ≥2 elite scalps is a **season-scale variance profile** — those scalps came from heat. When that team leads a structurally better opponent, the lead is exactly the unsustainable kind the system hunts. "Can beat anyone" and "can hold a lead" are opposite properties of the same coin.




Fig 6 — Inside the killer cell (pre-declared cuts only). Deficit stratifies hard — 77% at 1-4 down vs 52% at 5-9 — mirroring the global signal. Quarter is flat. Most notable: **live leader eFG heat adds nothing (65/65)** — the flag identifies variance at the season level, independent of whether the leader happens to be hot right now. Consistent with the locked "eFG-heat is not operative" finding.

SYSTEMATIC IDENTIFICATION — can we spot these at push time?

Yes — every input is as-of computable from data the nightly already maintains. No new sources, no look-ahead:

```
ELITE(date) = teams having reached win% >= .600 with >= 15 GP by date
KILLER(L) = leader has >= 2 as-of wins vs ELITE members [h2h table, stored]
BEATABLE(L) = leader as-of win% < .450 [team records, stored]
SPOT = trailer down 1-9, Q2/Q3, quality gap >= .15 [existing gates]
FIRE (ledger) = SPOT and BEATABLE(L) and KILLER(L) -> log, predicted ~65%
```

Decision-metric ladder (mirrors how GAP_BASE earned its way in)

Entry qualification — the three-condition flag above. It composes with, not replaces, the quality gap: the gap does the heavy lifting (56.4% base), the killer flag adds ~+8pp of separation on top.

Live confirmation — none required by the evidence. The one live overlay tested (leader eFG heat) added nothing; deficit ≤ 4 is the strongest secondary but is already how the band grades depth. The A-tier variance machinery stays A-tier's.

Sizing — nothing until the forward ledger earns it. Discretionary reads on killer-flagged leaders (POR, LA, CHI, SEA, PHX today; CON/TOR are never-beats-anyone, NOT flagged) belong in the bet log tagged for the discretionary-alpha stream.

CAUTION — power discipline: the star cell is $n=51$ (LOW). Direction is consistent across both seasons and mechanistically coherent, but ~65% is a hypothesis to be measured forward, not a deployable probability. That is precisely what the ledger stream is for. Multiplicity note: seven definitions were pre-registered; the killer result is the inversion of a pre-stated hypothesis, not a scan across cuts.

Dispositions & asks

Shipped/settled: def A stays production (gates, lane, registration). EVER600 adopted as descriptive lens. Production $< .400$ fade cut unchanged. **Open PM decisions:** (1) build KILLER_LEADER ledger stream — log-only, GAP_BASE-style, ~30 lines; (2) override-lane digest line — trigger fired (MIN $-.10 + LV -.05$ both lane-eligible per Jul 16 nightly); (3) follow-up cut on the same pool: killer-flagged *trailers* (the POR@CHI teaching-pair side) — is the variance profile equally disqualifying on the buy side?

Sources: research/2026-07-15_tier_definition_backtest.md ·
research/2026-07-16_beatable_leader_definitions.md · reproducible from export_xgb_training +
team_game_stats.